

Implementation of Quantum Superposition Gaming in Poker: Technical Architecture and Experimental Design

Srinivas
Lossfunk Research Laboratory

September 8, 2025

Abstract

This document provides comprehensive implementation details for Quantum Superposition Gaming (QSG) applied to Texas Hold'em poker. We present the complete technical architecture including quantum hand range representations, observation-collapse mechanisms, strategic deception algorithms, and training protocols. Our implementation extends classical poker AI by allowing agents to maintain probabilistic superpositions over hand ranges until strategic revelation events. We detail the mathematical formulations, neural network architectures, training algorithms, and experimental protocols necessary to validate quantum-inspired strategic deception in competitive poker environments.

1 Implementation Overview

1.1 System Architecture

The Quantum Poker Agent (QPA) consists of four primary components:

1. **Quantum Hand Range Module:** Maintains superposition states over possible hand holdings
2. **Observation Engine:** Detects collapse-inducing events from opponent actions
3. **Strategic Collapse Controller:** Determines optimal collapse timing and target states
4. **Quantum Policy Network:** Maps superposition states to betting actions

1.2 Environment Specifications

We implement QSG within a modified Texas Hold'em environment supporting:

- Standard 52-card deck with deterministic shuffling for reproducibility
- No-limit betting structure with configurable stack sizes
- 2-player heads-up format for simplified analysis
- Extended action space including quantum-specific operations
- Enhanced observation space capturing superposition states

2 Quantum Hand Range Representation

2.1 Mathematical Foundation

A quantum hand range is represented as a complex-valued vector over the space of all possible two-card combinations:

$$|\Psi_{\text{range}}\rangle = \sum_{h \in \mathcal{H}} \alpha_h e^{i\phi_h} |h\rangle \quad (1)$$

where:

- \mathcal{H} is the set of all 1326 possible hole card combinations
- $\alpha_h \in [0, 1]$ represents the amplitude for holding hand h
- $\phi_h \in [0, 2\pi]$ encodes strategic phase information
- Normalization constraint: $\sum_h |\alpha_h|^2 = 1$

2.2 Hand Range Evolution

Hand ranges evolve through unitary transformations based on community cards and strategic decisions:

$$|\Psi_{t+1}\rangle = U_{\text{community}}(c_t) \cdot U_{\text{strategy}}(\theta_t) \cdot |\Psi_t\rangle \quad (2)$$

where:

- $U_{\text{community}}(c_t)$ eliminates impossible hands given community card c_t
- $U_{\text{strategy}}(\theta_t)$ applies strategic amplitude modulation
- Both operators preserve normalization through unitary construction

2.3 Implementation Data Structures

```
class QuantumHandRange:
    def __init__(self, n_hands=1326):
        self.amplitudes = np.complex128(np.zeros(n_hands))
        self.phases = np.float64(np.zeros(n_hands))
        self.hand_map = self._generate_hand_mapping()

    def initialize_uniform_superposition(self):
        """Initialize equal superposition over all hands"""
        self.amplitudes = np.complex128(1.0/np.sqrt(1326))
        self.phases = np.random.uniform(0, 2*np.pi, 1326)

    def apply_community_card_constraint(self, community_cards):
        """Remove impossible hands given community cards"""
        valid_mask = self._get_valid_hand_mask(community_cards)
        self.amplitudes[~valid_mask] = 0.0
        self._renormalize()

    def get_collapsed_probabilities(self):
        """Return classical probabilities from quantum amplitudes"""
        return np.abs(self.amplitudes)**2
```

3 Observation and Collapse Mechanisms

3.1 Observation Events

We define specific opponent actions that constitute observation events forcing superposition collapse:

1. **Aggressive Betting:** Bets $> 2.5 \times$ pot size
2. **All-in Actions:** Complete stack commitment
3. **Unusual Sizing:** Bet sizes outside normal ranges
4. **Temporal Patterns:** Rapid succession of actions
5. **Showdown:** Mandatory full collapse event

3.2 Collapse Probability Computation

The probability of collapse given opponent action a_{opp} is computed as:

$$P(\text{collapse}|a_{\text{opp}}) = \sigma(\mathbf{w}^T \phi(a_{\text{opp}}, |\Psi\rangle, \text{context})) \quad (3)$$

where:

- $\phi(\cdot)$ extracts features from action, superposition state, and game context
- \mathbf{w} are learned collapse sensitivity parameters
- $\sigma(\cdot)$ is the sigmoid function ensuring $P \in [0, 1]$

3.3 Strategic Collapse Selection

When collapse occurs, the agent strategically selects the target collapsed state:

$$h_{\text{collapse}} = \arg \max_h Q_{\text{collapse}}(h||\Psi\rangle, a_{\text{opp}}, s_t) \quad (4)$$

where Q_{collapse} is a learned value function estimating the strategic benefit of collapsing to specific hands.

3.4 Collapse Implementation

```
class CollapseEngine:
    def __init__(self, collapse_net, strategy_net):
        self.collapse_net = collapse_net  # Neural network
        self.strategy_net = strategy_net

    def should_collapse(self, opponent_action, quantum_range, context):
        """Determine if opponent action triggers collapse"""
        features = self._extract_collapse_features(
            opponent_action, quantum_range, context
        )
        collapse_prob = torch.sigmoid(self.collapse_net(features))
        return np.random.random() < collapse_prob.item()

    def select_collapse_target(self, quantum_range, opponent_action,
                              context):
        """Choose which hand to collapse to"""
```

```

valid_hands = quantum_range.get_nonzero_hands()
collapse_values = []

for hand in valid_hands:
    hand_features = self._encode_hand_context(
        hand, opponent_action, context
    )
    value = self.strategy_net(hand_features)
    collapse_values.append(value.item())

best_hand_idx = np.argmax(collapse_values)
return valid_hands[best_hand_idx]

```

4 Neural Network Architecture

4.1 Quantum Policy Network

The quantum policy network processes superposition states and outputs betting actions:

$$\text{QPN}(|\Psi\rangle, s_t) = \text{softmax}(\text{MLP}([\phi_{\text{quantum}}, \phi_{\text{game}}])) \quad (5)$$

where:

- $\phi_{\text{quantum}} = [\text{Re}(\alpha_h), \text{Im}(\alpha_h), \phi_h]_{h \in \mathcal{H}}$ encodes superposition
- ϕ_{game} contains standard poker features (pot size, position, etc.)
- MLP uses complex-valued activations to preserve quantum information

4.2 Complex-Valued Neural Layers

We implement custom neural layers for complex-valued computations:

```

class ComplexLinear(nn.Module):
    def __init__(self, in_features, out_features):
        super().__init__()
        self.real_weight = nn.Parameter(torch.randn(out_features,
            in_features))
        self.imag_weight = nn.Parameter(torch.randn(out_features,
            in_features))
        self.real_bias = nn.Parameter(torch.randn(out_features))
        self.imag_bias = nn.Parameter(torch.randn(out_features))

    def forward(self, x_complex):
        x_real, x_imag = x_complex.real, x_complex.imag

        out_real = (torch.mm(x_real, self.real_weight.t()) -
            torch.mm(x_imag, self.imag_weight.t()) +
            self.real_bias)
        out_imag = (torch.mm(x_real, self.imag_weight.t()) +
            torch.mm(x_imag, self.real_weight.t()) +
            self.imag_bias)

        return torch.complex(out_real, out_imag)

class QuantumPolicyNetwork(nn.Module):

```

```

def __init__(self, n_hands=1326, hidden_dim=512):
    super().__init__()
    self.quantum_encoder = ComplexLinear(n_hands * 3, hidden_dim)
    self.game_encoder = nn.Linear(64, hidden_dim) # Standard
    features
    self.fusion = ComplexLinear(hidden_dim, hidden_dim)
    self.output = nn.Linear(hidden_dim, 4) # fold, call, raise,
    quantum_action

def forward(self, quantum_range, game_state):
    # Encode quantum superposition
    quantum_features = torch.cat([
        quantum_range.amplitudes.real,
        quantum_range.amplitudes.imag,
        quantum_range.phases
    ])
    quantum_encoded = self.quantum_encoder(quantum_features)

    # Encode standard game features
    game_encoded = self.game_encoder(game_state)
    game_complex = torch.complex(game_encoded, torch.zeros_like(
        game_encoded))

    # Fuse quantum and classical information
    fused = self.fusion(quantum_encoded + game_complex)
    action_logits = self.output(fused.real) # Project to real for
    actions

    return torch.softmax(action_logits, dim=-1)

```

4.3 Value Function Architecture

The quantum value function estimates expected returns from superposition states:

$$V^\pi(|\Psi\rangle, s_t) = \mathbb{E}_{h \sim |\Psi\rangle} \left[\sum_{\tau=t}^T \gamma^{\tau-t} r_\tau \right] \quad (6)$$

Implementation uses Monte Carlo sampling over possible collapse trajectories:

```

class QuantumValueFunction(nn.Module):
    def __init__(self, n_hands=1326, hidden_dim=256):
        super().__init__()
        self.quantum_processor = ComplexLinear(n_hands * 3, hidden_dim)
        self.game_processor = nn.Linear(64, hidden_dim)
        self.value_head = nn.Linear(hidden_dim * 2, 1)

    def forward(self, quantum_range, game_state):
        q_features = self._encode_quantum_state(quantum_range)
        q_processed = self.quantum_processor(q_features).real

        g_processed = self.game_processor(game_state)

        combined = torch.cat([q_processed, g_processed], dim=-1)
        return self.value_head(combined)

    def monte_carlo_evaluation(self, quantum_range, game_state,
        n_samples=100):

```

```

"""Evaluate superposition through sampling"""
values = []
for _ in range(n_samples):
    # Sample hand from superposition
    hand = quantum_range.sample_hand()
    # Evaluate classical value with sampled hand
    classical_state = self._create_classical_state(hand,
        game_state)
    value = self.forward_classical(classical_state)
    values.append(value)
return torch.mean(torch.stack(values))

```

5 Training Algorithm

5.1 Quantum Policy Gradient

We extend standard policy gradient methods to handle superposition states:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{|\Psi\rangle \sim \pi_{\theta}} \left[\sum_t \nabla_{\theta} \log \pi_{\theta}(a_t || \Psi_t) \cdot A_t \right] \quad (7)$$

where the advantage function A_t accounts for quantum uncertainty:

$$A_t = Q^{\pi}(|\Psi_t\rangle, a_t) - V^{\pi}(|\Psi_t\rangle) \quad (8)$$

5.2 Collapse Strategy Learning

The collapse strategy is trained using a separate objective that maximizes expected value improvement:

$$\mathcal{L}_{\text{collapse}} = \mathbb{E} \left[(V^{\pi}(h_{\text{collapse}}) - V^{\pi}(|\Psi\rangle))^2 \right] \quad (9)$$

This encourages the agent to collapse to hands that improve its strategic position.

5.3 Deception Reward Shaping

We introduce additional reward terms that encourage strategic deception:

$$r_{\text{total}} = r_{\text{game}} + \lambda_1 r_{\text{superposition}} + \lambda_2 r_{\text{deception}} + \lambda_3 r_{\text{collapse}} \quad (10)$$

where:

- $r_{\text{superposition}} = \alpha \cdot \text{entropy}(|\Psi\rangle)$ rewards maintaining uncertainty
- $r_{\text{deception}} = \beta \cdot \text{KL}(\hat{P}_{\text{opp}} || P_{\text{true}})$ rewards misleading opponent beliefs
- $r_{\text{collapse}} = \gamma \cdot \mathbb{I}[\text{strategic collapse}]$ rewards well-timed revelation

5.4 Training Loop Implementation

```
class QuantumPokerTrainer:
    def __init__(self, policy_net, value_net, collapse_net):
        self.policy_net = policy_net
        self.value_net = value_net
        self.collapse_net = collapse_net
        self.optimizer = torch.optim.Adam(self.get_all_parameters())

    def train_episode(self, opponent_agent):
        episode_data = []
        quantum_range = QuantumHandRange()
        quantum_range.initialize_from_dealt_cards(self.deal_cards())

        game_state = self.env.reset()
        done = False

        while not done:
            # Get action from quantum policy
            action_probs = self.policy_net(quantum_range, game_state)
            action = self.sample_action(action_probs)

            # Execute action and observe
            next_state, reward, done, info = self.env.step(action)

            # Check for collapse events
            opponent_action = info.get('opponent_action')
            if self.collapse_engine.should_collapse(opponent_action,
                                                    quantum_range):
                collapsed_hand = self.collapse_engine.
                    select_collapse_target(
                        quantum_range, opponent_action, game_state
                    )
                quantum_range.collapse_to_hand(collapsed_hand)

            # Store transition
            episode_data.append({
                'quantum_state': quantum_range.copy(),
                'game_state': game_state,
                'action': action,
                'reward': reward,
                'next_state': next_state
            })

            game_state = next_state

        # Update networks using collected data
        self.update_networks(episode_data)

    def update_networks(self, episode_data):
        policy_loss = self.compute_policy_loss(episode_data)
        value_loss = self.compute_value_loss(episode_data)
        collapse_loss = self.compute_collapse_loss(episode_data)

        total_loss = policy_loss + value_loss + collapse_loss

        self.optimizer.zero_grad()
        total_loss.backward()
```

6 Experimental Protocol

6.1 Training Setup

Environment Configuration:

- Starting stack: 200 big blinds
- Blind structure: 1/2 fixed throughout training
- Hand history logging for detailed analysis
- Deterministic opponent for initial training

Network Hyperparameters:

- Learning rate: $\alpha = 3 \times 10^{-4}$ with cosine annealing
- Batch size: 64 episodes
- Hidden dimensions: 512 for policy, 256 for value
- Gradient clipping: max norm 0.5
- Discount factor: $\gamma = 0.99$

Quantum-Specific Parameters:

- Superposition reward weight: $\lambda_1 = 0.1$
- Deception reward weight: $\lambda_2 = 0.05$
- Collapse reward weight: $\lambda_3 = 0.02$

6.2 Opponent Models

We train against progressively sophisticated opponents:

1. **Random Agent:** Baseline for sanity checking
2. **Calling Station:** Always calls, never raises
3. **Tight-Aggressive Bot:** Classical optimal strategy approximation
4. **Libratus-Style Agent:** State-of-the-art classical poker AI
5. **Human Experts:** Professional poker players for validation

6.3 Evaluation Metrics

Performance Metrics:

- Expected value per hand (bb/100)
- Win rate across different opponent types
- Variance analysis for risk assessment
- Convergence speed to optimal policies

Quantum-Specific Metrics:

- Superposition maintenance duration
- Collapse timing optimality
- Deception success rate (measured via opponent belief tracking)
- Quantum strategy diversity (entropy of discovered patterns)

Emergent Behavior Analysis:

- Novel betting patterns not seen in classical poker
- Strategic collapse timing patterns
- Quantum bluffing frequency and success
- Opponent confusion metrics

7 Implementation Challenges and Solutions

7.1 Computational Complexity

Challenge: Maintaining superposition over 1326 possible hands is computationally expensive.

Solution: Implement sparse superposition representation focusing on strategically relevant hands:

```
class SparseQuantumRange:
    def __init__(self, sparsity_threshold=1e-6):
        self.active_hands = {} # hand_id -> amplitude
        self.threshold = sparsity_threshold

    def prune_negligible_amplitudes(self):
        """Remove hands with negligible probability"""
        to_remove = [h for h, amp in self.active_hands.items()
                     if abs(amp)**2 < self.threshold]
        for h in to_remove:
            del self.active_hands[h]

    def adaptive_sparsity(self, target_size=100):
        """Keep only top-k most probable hands"""
        if len(self.active_hands) <= target_size:
            return
        sorted_hands = sorted(self.active_hands.items(),
                              key=lambda x: abs(x[1])**2, reverse=True)
        self.active_hands = dict(sorted_hands[:target_size])
        self._renormalize()
```

7.2 Training Stability

Challenge: Complex-valued networks can exhibit training instabilities.

Solution: Implement specialized initialization and regularization:

```
def initialize_complex_weights(layer):
    """Initialize complex weights for stable training"""
    if isinstance(layer, ComplexLinear):
        # Xavier initialization for complex weights
        fan_in = layer.real_weight.size(1)
        std = np.sqrt(2.0 / fan_in)
        layer.real_weight.data.normal_(0, std)
        layer.imag_weight.data.normal_(0, std)
        layer.real_bias.data.zero_()
        layer.imag_bias.data.zero_()

def complex_gradient_clipping(model, max_norm=1.0):
    """Clip gradients for complex-valued parameters"""
    total_norm = 0
    for p in model.parameters():
        if p.grad is not None:
            param_norm = p.grad.data.norm(2)
            total_norm += param_norm.item() ** 2
    total_norm = total_norm ** (1. / 2)

    clip_coef = max_norm / (total_norm + 1e-6)
    if clip_coef < 1:
        for p in model.parameters():
            if p.grad is not None:
                p.grad.data.mul_(clip_coef)
```

7.3 Opponent Belief Modeling

Challenge: Measuring deception success requires accurate opponent belief tracking.

Solution: Implement Bayesian opponent model with uncertainty quantification:

```
class OpponentBeliefTracker:
    def __init__(self, n_hands=1326):
        self.belief_history = []
        self.action_history = []
        self.hand_prior = np.ones(n_hands) / n_hands

    def update_belief(self, opponent_action, community_cards):
        """Update belief about opponent hand range"""
        # Bayes update based on action
        likelihood = self._compute_action_likelihood(
            opponent_action, community_cards
        )

        posterior = self.hand_prior * likelihood
        posterior /= np.sum(posterior)

        self.belief_history.append(posterior.copy())
        self.action_history.append(opponent_action)
        self.hand_prior = posterior

    def compute_deception_metric(self, true_quantum_range):
        """Measure how much opponent belief differs from reality"""
```

```

if not self.belief_history:
    return 0.0

true_probs = true_quantum_range.get_collapsed_probabilities()
believed_probs = self.belief_history[-1]

# KL divergence as deception measure
kl_div = np.sum(true_probs * np.log(
    (true_probs + 1e-8) / (believed_probs + 1e-8)
))
return kl_div

```

8 Expected Results and Analysis

8.1 Predicted Emergent Behaviors

Quantum Bluffing Patterns:

- Maintaining superposition over strong and weak hands simultaneously
- Strategic collapse to unexpected hand strengths
- Temporal deception through delayed revelation

Novel Betting Strategies:

- Bet sizes that reflect superposition uncertainty
- Action sequences impossible in classical poker
- Meta-strategic collapse timing

Opponent Exploitation:

- Learning opponent collapse triggers
- Exploiting opponent belief models
- Creating cognitive overload through quantum complexity

8.2 Performance Predictions

Based on theoretical analysis, we expect:

- 15-25% improvement in expected value against classical opponents
- Novel strategic patterns not present in existing poker literature
- Emergent quantum strategies that transfer to other game domains
- Demonstration of genuine strategic advantages from superposition maintenance

9 Conclusion

This implementation guide provides a complete technical roadmap for creating the first quantum superposition gaming agent applied to poker. The combination of rigorous mathematical foundations, practical implementation details, and comprehensive experimental protocols creates a framework for validating quantum-inspired strategic deception in competitive environments.

The poker domain serves as an ideal testbed for QSG principles, offering natural hidden information structures, clear observation mechanics, and measurable strategic outcomes. Success in this implementation would establish the viability of quantum superposition concepts for broader game-theoretic applications while opening entirely new research directions in strategic AI.

The technical challenges outlined here, while significant, are addressable through the proposed solutions and represent genuine contributions to both reinforcement learning and game theory. The expected emergent behaviors would constitute novel strategic patterns previously impossible in classical gaming frameworks.

References